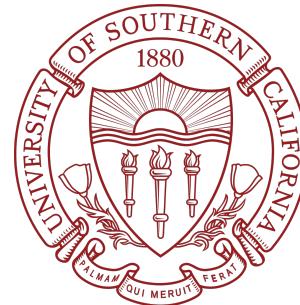


Graph Learning for Bipartite Networks



Tian Xie

Collaboration with Chaoyang He, Xiang Ren, Cyrus Shahabi,
and C.-C. Jay Kuo

Real-world Networks



Academic Graph



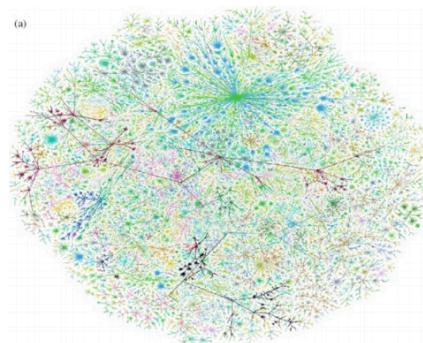
Office/Social Graph



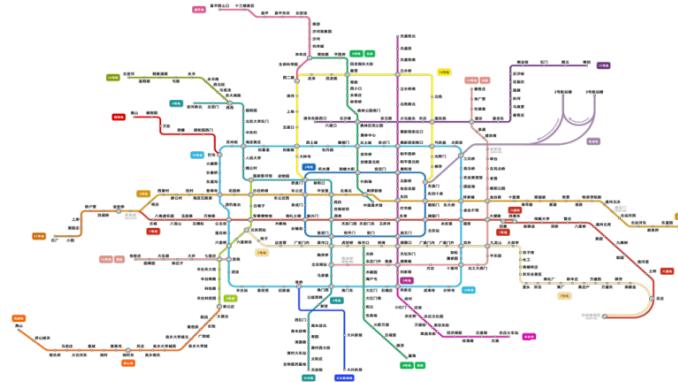
Biological Neural Networks



Knowledge Graph

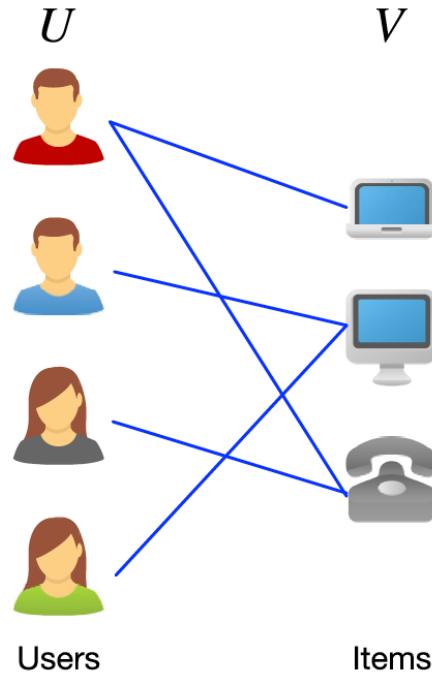


Internet



Transportation

Bipartite Graphs

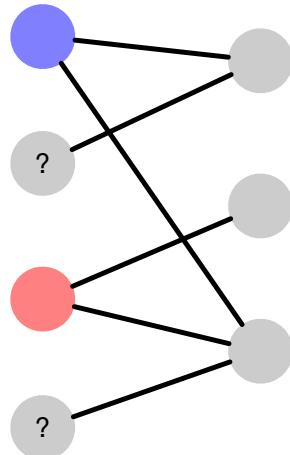


Examples:

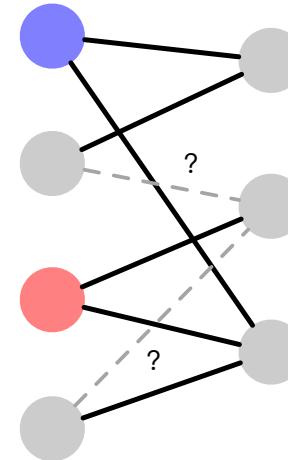
- **E-commerce**: Users – Items;
- **Online ratings**: Users – Movies;
- **Academic**: Authors – Papers;
- **Search engines**: Queries – Answers;
- ...

Machine Learning Tasks on Graphs

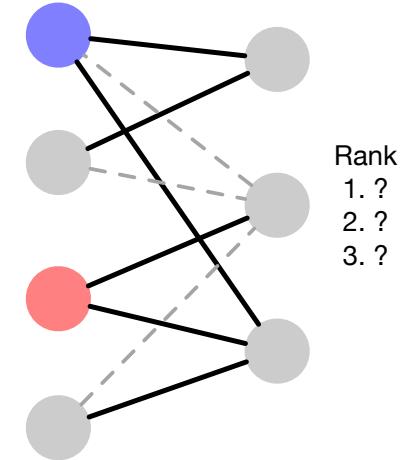
What problems are we trying to solve?



Node classification



Link prediction



Recommendation

More: community detection, anomaly detection, network similarity...

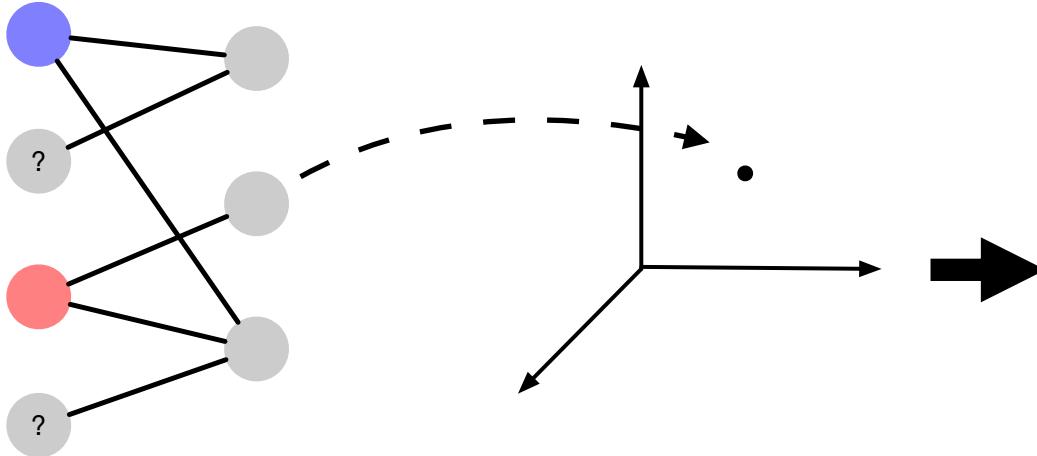
Outline

- Background & Related Work
- Challenges for Bipartite Graphs
- Our Method: L-BGNN
- Experiments & Evaluation
- Conclusion

Outline

- Background & Related Work
 - Graph Representation Learning
- Challenges for Bipartite Graphs
- Our Method: L-BGNN
- Experiments & Evaluation
- Conclusion

Traditional Learning Paradigm



- 1. Node attributes;
- 2. Network structure;

Hand-crafted features:

- 1. Node attributes;
- 2. Network structure: statistic of the network
 - Node degree;
 - Eccentricity;
 - Radius;
 - ...

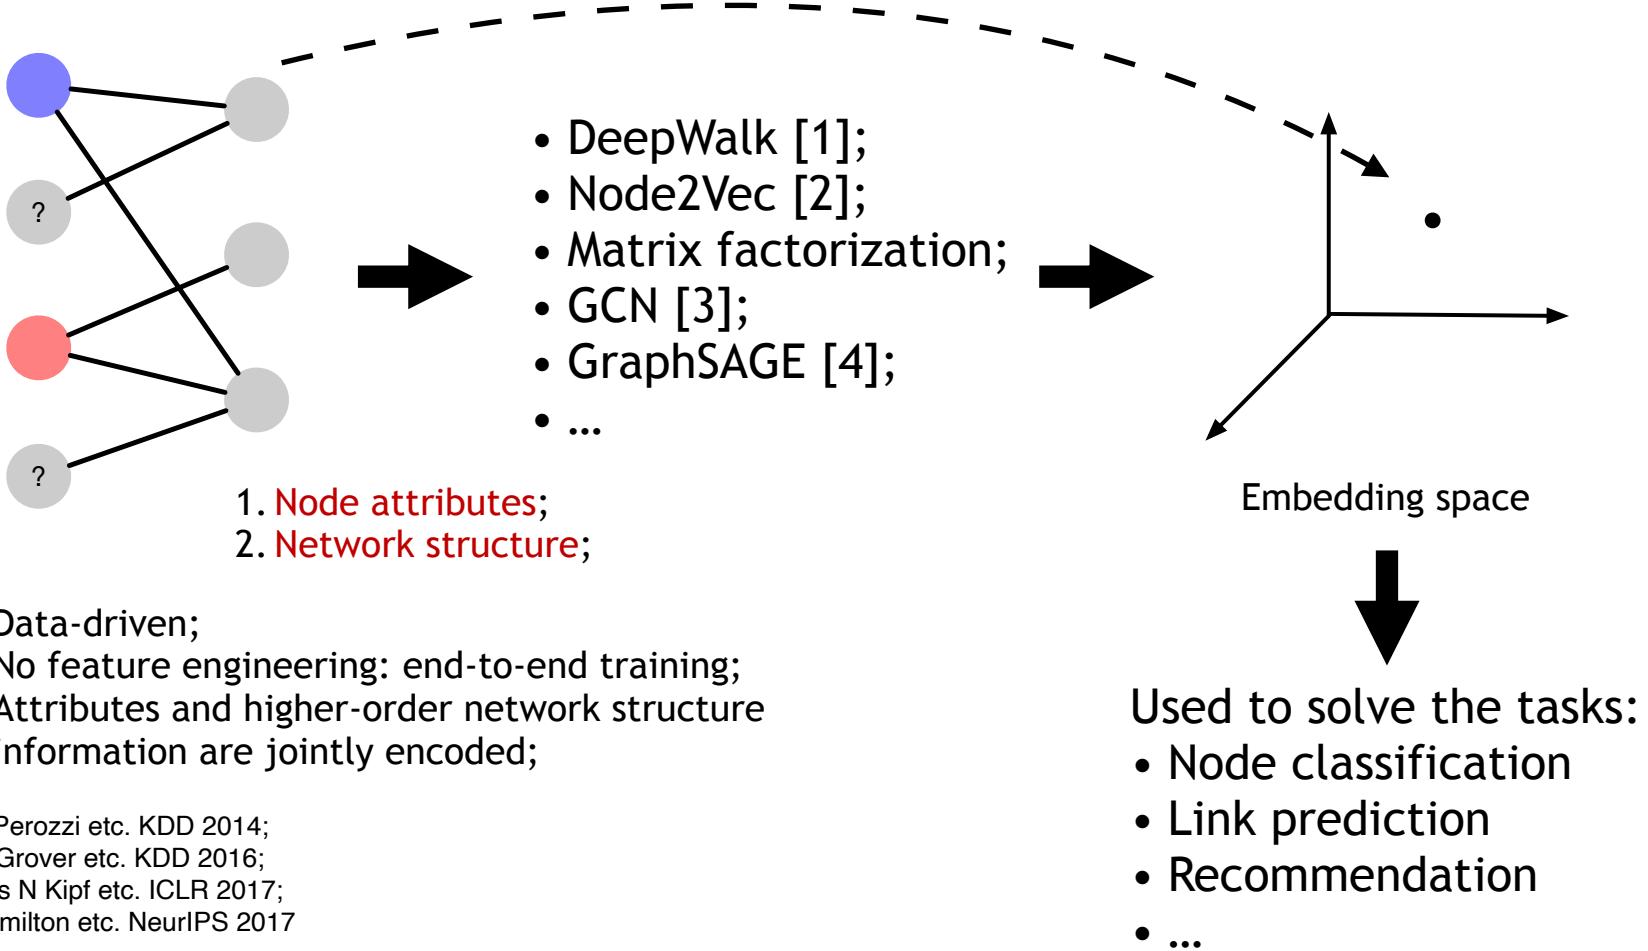
Used to solve the tasks:

- Node classification
- Link prediction
- Recommendation
- ...

Machine learning models

- Logistic regression;
- Random forest;
- Neural networks;
- ...

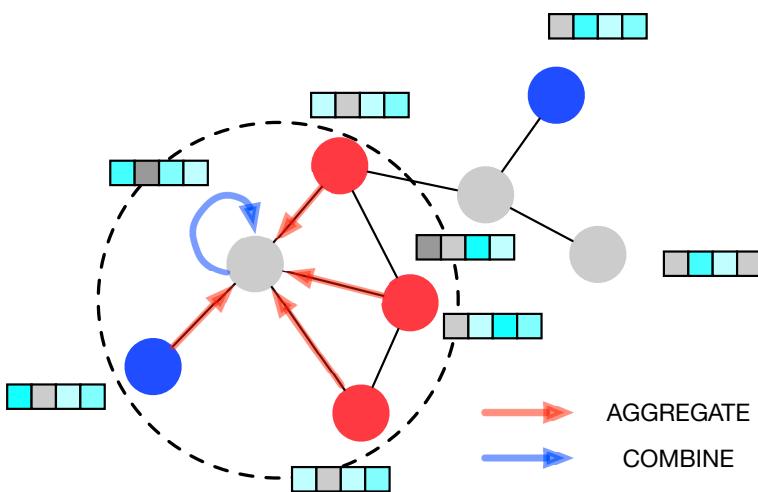
Graph Representation Learning



Outline

- Background & Related Work
 - Graph Representation Learning
 - Graph Neural Networks
- Challenges for Bipartite Graphs
- Our Method: L-BGNN
- Experiments & Evaluation
- Conclusion

Graph Neural Networks



- Message passing: a general framework

$$a_v^{(k)} = \text{AGGREGATE}^{(k)}\left(\left\{h_u^{(k-1)} : u \in \mathcal{N}(v)\right\}\right)$$

$$h_v^{(k)} = \text{COMBINE}^{(k)}\left(h_v^{(k-1)}, a_v^{(k)}\right)$$

$$\mathcal{L} = Loss(h^{(K)}, y)$$

Layer

- Graph Convolutional Network (GCN)

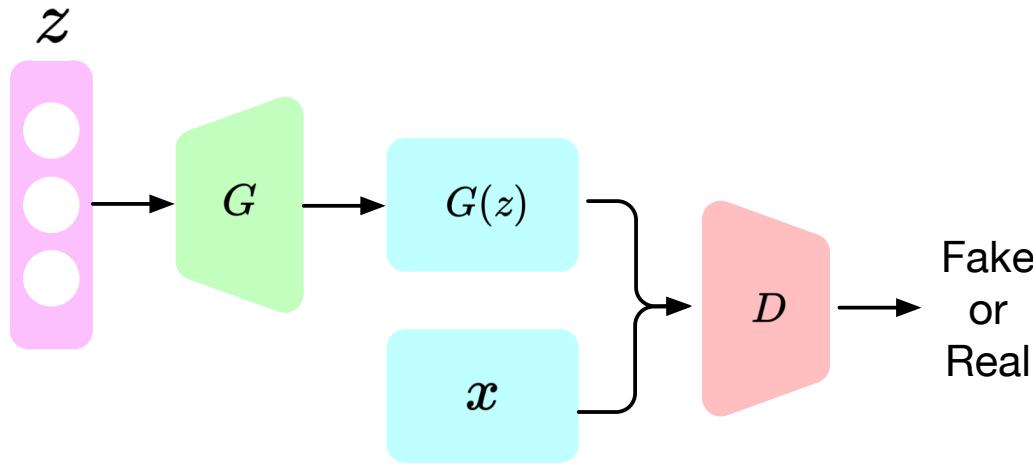
$$h_v^{(k)} = \text{ReLU}\left(W \cdot \text{MEAN}\left\{h_u^{(k-1)}, \forall u \in \mathcal{N}(v) \cup \{v\}\right\}\right)$$

$$\mathcal{L} = H(h^{(K)}, y)$$

Layer

[1] Thomas N Kipf etc. ICLR 2017;
[2] Keyulu Xu etc. ICLR 2019;

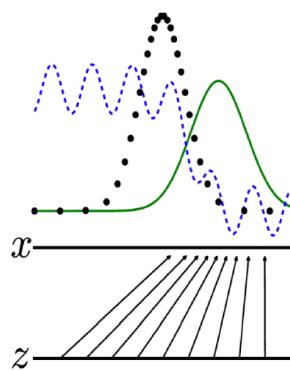
Generative Adversarial Networks (GANs)



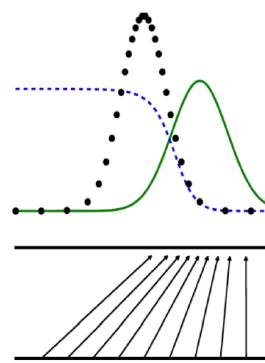
$$\min_{\theta^G} \max_{\theta^D} \mathbb{E}_{x \sim p_{\text{data}}} [\log D(x; \theta^D)] + \mathbb{E}_{z \sim p_g} [\log(1 - D(G(z; \theta^G); \theta^D))]$$

[1] Ian Goodfellow etc. NeurIPS 2014.

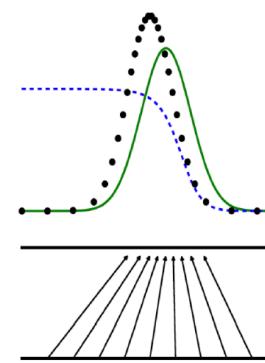
- Non-linear distributions mapping



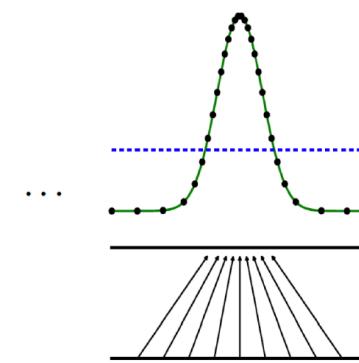
(a)



(b)



(c)

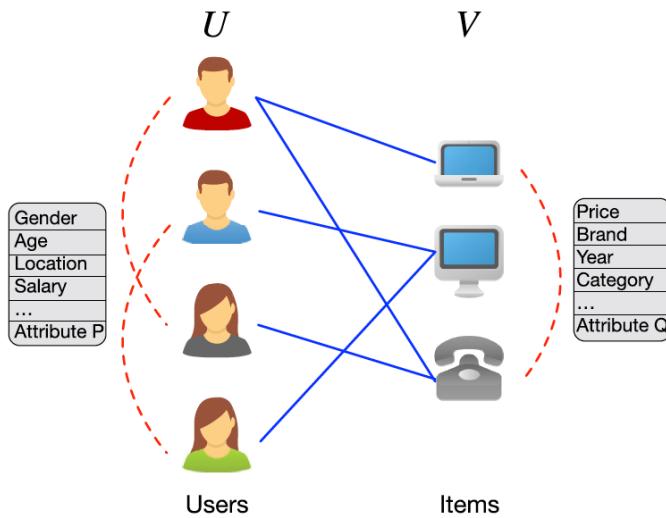


(d)

Outline

- Background & Related Work
- Challenges for Bipartite Graphs
 - Domain Inconsistency
 - Efficiency
- Our Method: L-BGNN
- Experiments & Evaluation
- Conclusion

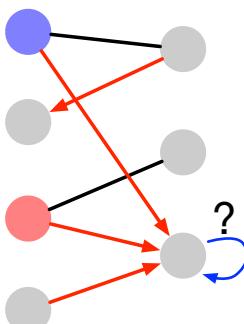
Domain Inconsistency



- Distinct attribute distributions;
- Message passing is **failed** or **suboptimal**;

$$a_v^{(k)} = \text{AGGREGATE}^{(k)} \left(\left\{ h_u^{(k-1)} : u \in \mathcal{N}(v) \right\} \right)$$

$$h_v^{(k)} = \text{COMBINE}^{(k)} \left(h_v^{(k-1)}, a_v^{(k)} \right)$$



Possible solutions:

- Convert to graphs with one domain?
 - **No feature correlations**
- Projection or concatenation?
 - **Too simple**
- Heterogeneous graph?
 - **Only one type of meta-path**

Efficiency

- E-commerce networks: millions of nodes and edges

- • Scalable & Label Efficiency
 - GNNs are hard for scaling;
 - Extremely few labeled samples (unsupervised);

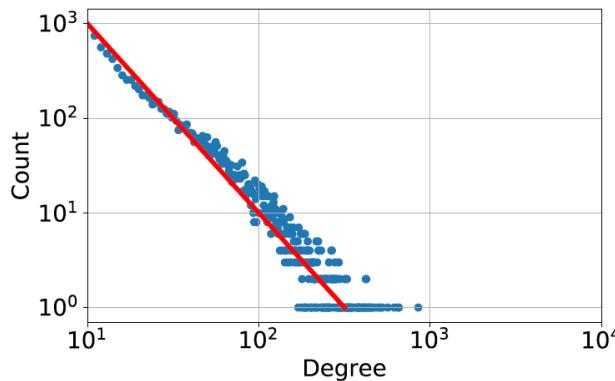


Fig. 2. Power-law distribution of node degrees in a user–item network from the Amazon dataset [33].

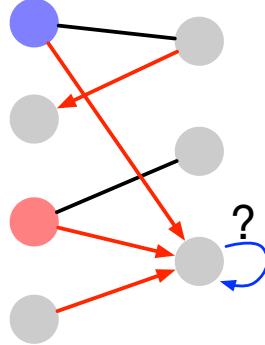
Possible solution:

- Sampling nodes or subgraphs?
 - High bias and variance for the message passing

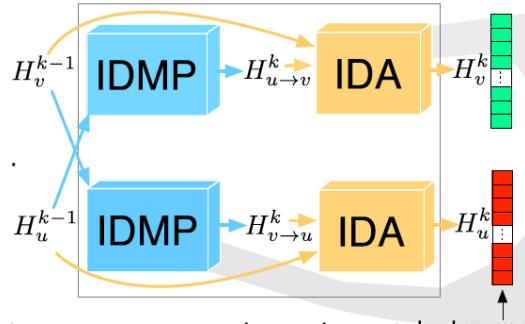
Outline

- Background & Related Work
- Challenges for Bipartite Graphs
- Our Method: L-BGNN
- Experiments & Evaluation
- Conclusion

L-BGNN: Layerwise Trained Bipartite Graph Neural Network

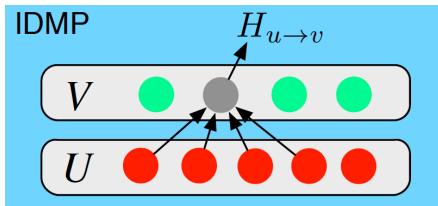


- First challenge: domain inconsistency
- Proposed solutions
 - Interdomain message passing (IDMP)
 - Intradomain alignment (IDA)



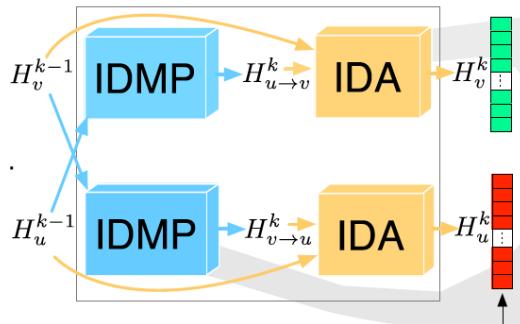
IDMP \approx Aggregate operator
IDA \approx Combine operator

IDMP: Interdomain Message Passing

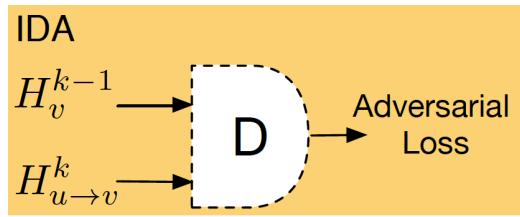
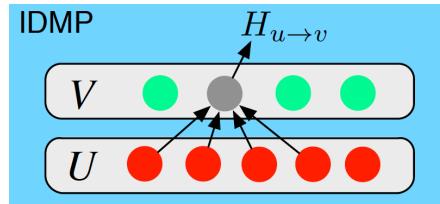


- Aggregate message;
- Independent filter parameters;

$$\begin{cases} \mathbf{H}_{v \rightarrow u}^{(k)} = \sigma(\hat{\mathbf{B}}_u \mathbf{H}_v^{(k)} \mathbf{W}_u^{(k)}) \\ \mathbf{H}_{u \rightarrow v}^{(k)} = \sigma(\hat{\mathbf{B}}_v \mathbf{H}_u^{(k)} \mathbf{W}_v^{(k)}) \end{cases} \quad \mathbf{A} = \begin{pmatrix} \mathbf{0}_u & \mathbf{B}_u \\ \mathbf{B}_v & \mathbf{0}_v \end{pmatrix}$$



IDA: Intradomain Alignment



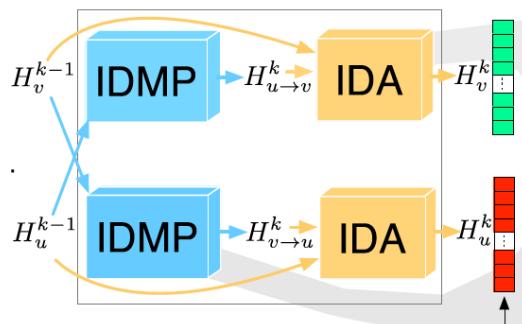
- Map domain distributions in an adversarial way.
- Intuition: neighbor nodes should share similar embedding distributions.

IDA (discriminator) loss:

$$\begin{aligned}\mathcal{L}_D(\theta^D | \theta^G) = & -\frac{1}{M} \sum_{i=1}^M \log P_{\theta^D, \theta^G}(\text{source} = 0 | \mathbf{h}_{u,i}^{(k)}) \\ & -\frac{1}{M} \sum_{i=1}^M \log P_{\theta^D, \theta^G}(\text{source} = 1 | \mathbf{h}_{v \rightarrow u,i}^{(k)})\end{aligned}$$

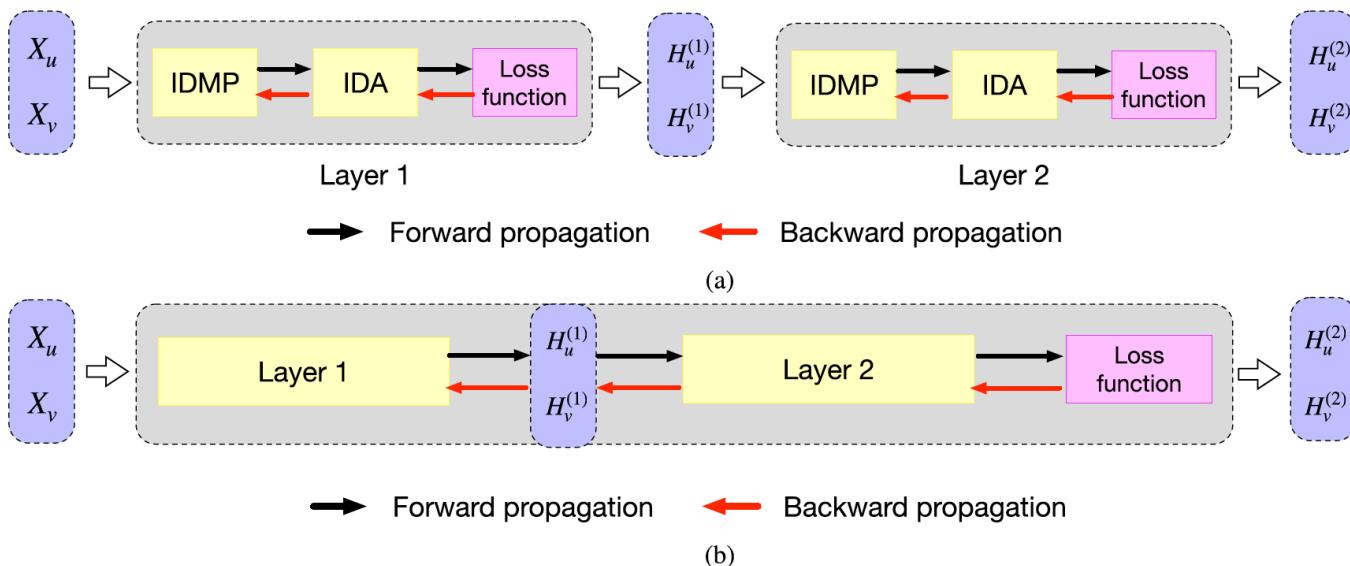
IDMP (generator) loss:

$$\mathcal{L}_G(\theta^G | \theta^D) = -\frac{1}{M} \sum_{i=1}^M \log P_{\theta^D, \theta^G}(\text{source} = 0 | \mathbf{h}_{v \rightarrow u,i}^{(k)})$$

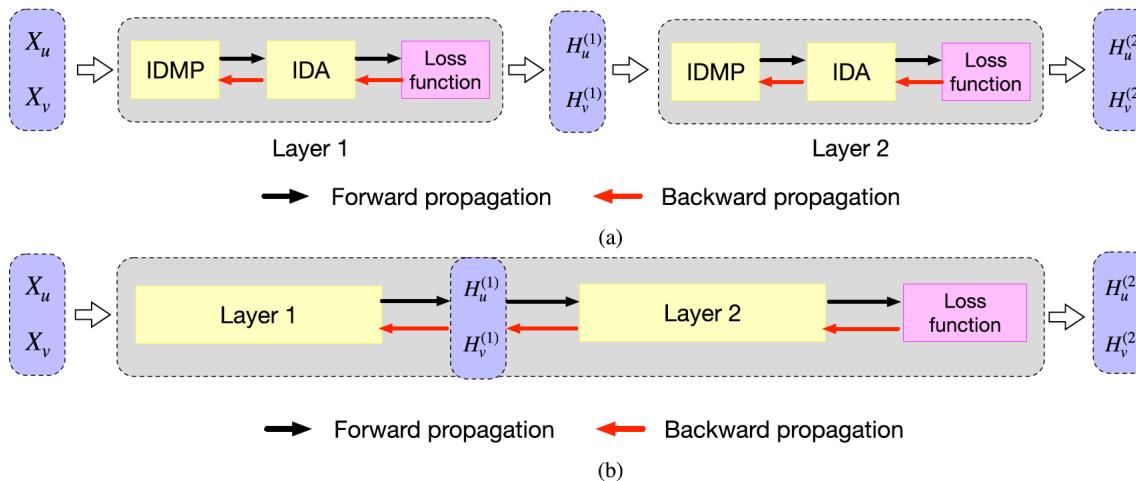


- Second challenge: scalable & label efficiency
- **Proposed solutions:** Layerwise training
 - Predictions from the IDMP are served as input in next training layer.

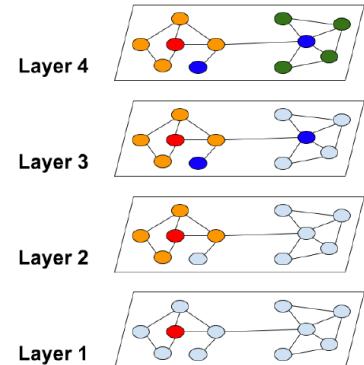
$$\begin{cases} \mathbf{H}_u^{(k+1)} = \mathbf{H}_{v \rightarrow u}^{(k)} \\ \mathbf{H}_v^{(k+1)} = \mathbf{H}_{u \rightarrow v}^{(k)} \end{cases}$$



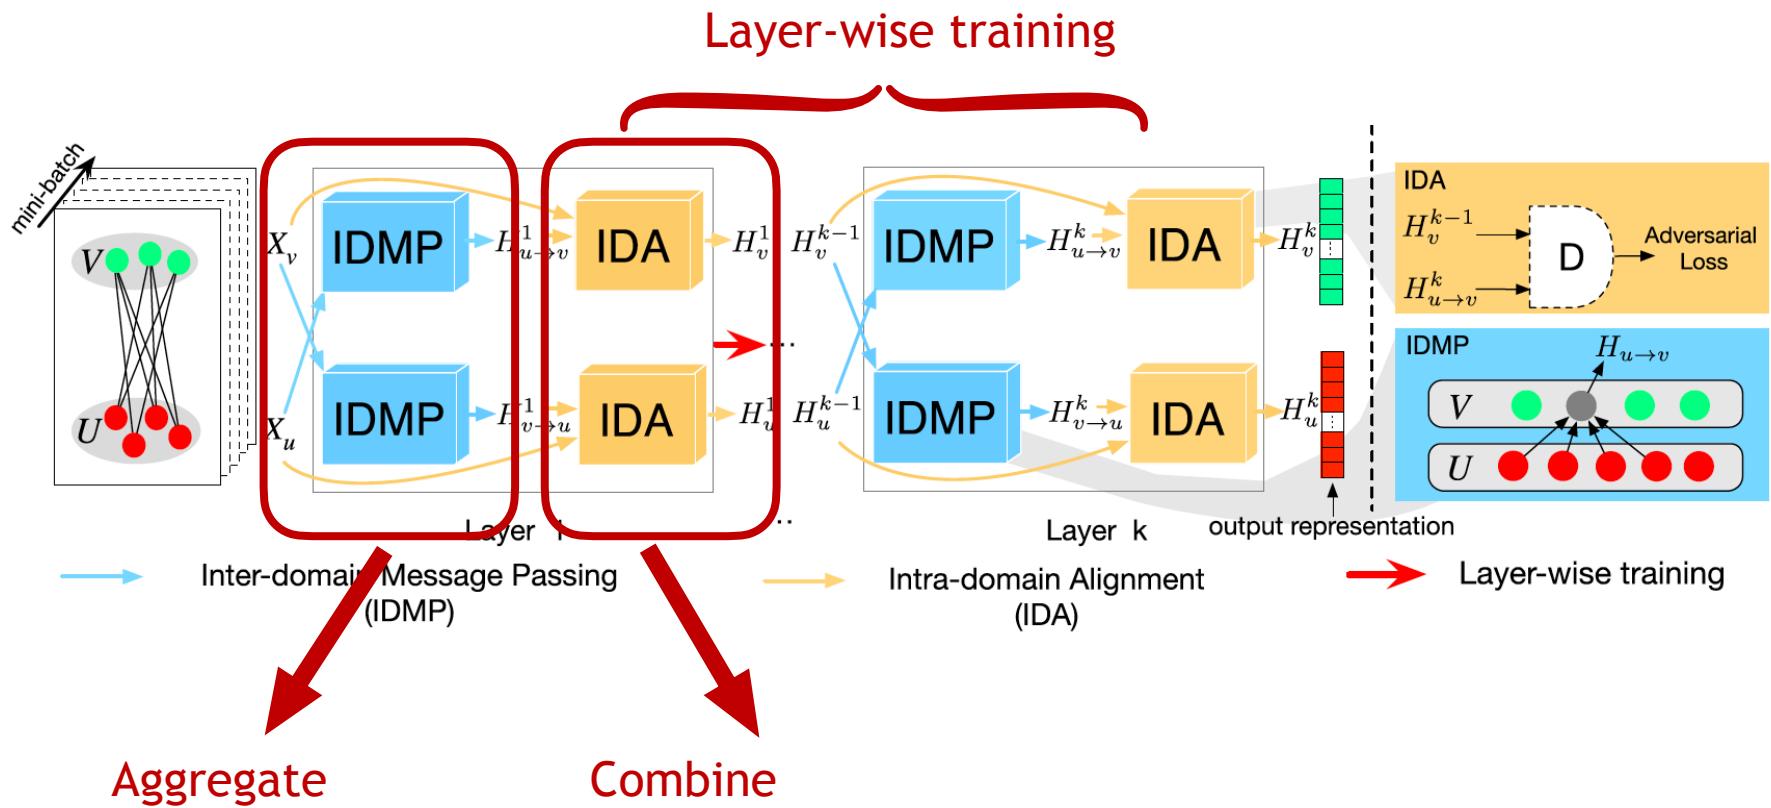
Layerwise Training



- Scalable
 - Only one training layer is alive;
 - No exponential neighborhood expansion: minibatch training;
- Without statistical performance sacrifice (Theorem 1)
 - Also can encode higher-order information;



L-BGNN



Outline

- Background & Related Work
- Challenges for Bipartite Graphs
- Our Method: L-BGNN
- Experiments & Evaluation
- Conclusion

Datasets

TABLE I
STATISTICS OF EIGHT DATASETS USED IN OUR EXPERIMENTS

Datasets	#Nodes U	#Nodes V	#Edges	#Feat U	#Feat V	#Classes
Group	619,030	90,044	991,734	8	16	2
Cora	734	877	1,802	1,433	1,000	7
CiteSeer	613	510	1,000	3,703	3,000	6
PubMed	13,424	3,435	18,782	400	500	3
Movielens	6,040	3,706	1,000,209	N.A.	N.A.	
DBLP	6,001	1,308	29,254	N.A.	N.A.	
Yelp	2,614	1,286	30,837	N.A.	N.A.	
Amazon	6,170	2,753	195,790	N.A.	N.A.	

- Real datasets

- Movielens (user – movies)
- DBLP (author – paper)
- Yelp (user – business)
- Amazon (user – item)
- Group (user – community)

- Synthetic datasets

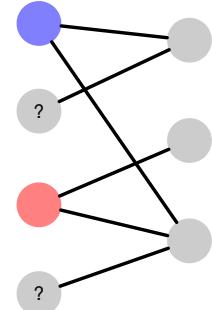
- Cora
- CiteSeer
- PubMed

Evaluation: Node Classification

TABLE II

PERFORMANCE COMPARISON FOR THE NODE CLASSIFICATION TASK, WHERE OOM DENOTES “OUT OF MEMORY”

Methods	Group	Cora		Citeseer		PubMed	
		Accuracy	Micro F1	Macro F1	Micro F1	Macro F1	Micro F1
node2vec		0.577	0.682	0.663	0.569	0.476	0.700
ANRL		OOM	0.797	0.785	0.748	0.678	0.390
VGAE		OOM	0.782	0.754	0.732	0.645	0.823
GraphSAGE-MEAN		0.580	0.823	0.801	0.748	0.665	0.843
HeGAN	> 2 days		0.189	0.167	0.228	0.224	0.396
FeatWalk		0.514	0.237	0.164	0.276	0.213	0.408
metapath2vec		OOM	0.628	0.596	0.577	0.466	0.609
HIN2Vec		0.518	0.635	0.593	0.561	0.514	0.630
BiNE	> 2 days		0.372	0.077	0.293	0.186	0.391
BiANE	> 2 days		0.791	0.773	0.739	0.654	0.793
L-BGNN-Adv		0.622	0.859	0.831	0.768	0.698	0.857
							0.860

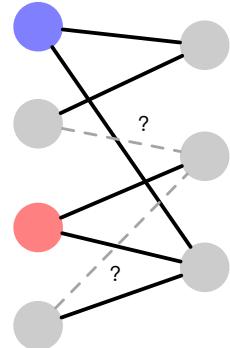


- Train a logistic classifier on the node embeddings;
- Metrics: accuracy; micro (macro) F1;
- SOTA performance, especially on the large-scale graph;

Evaluation: Link Prediction

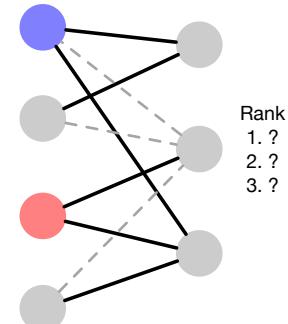
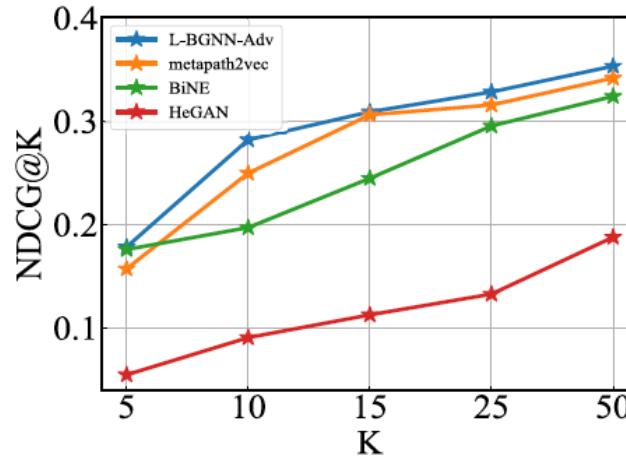
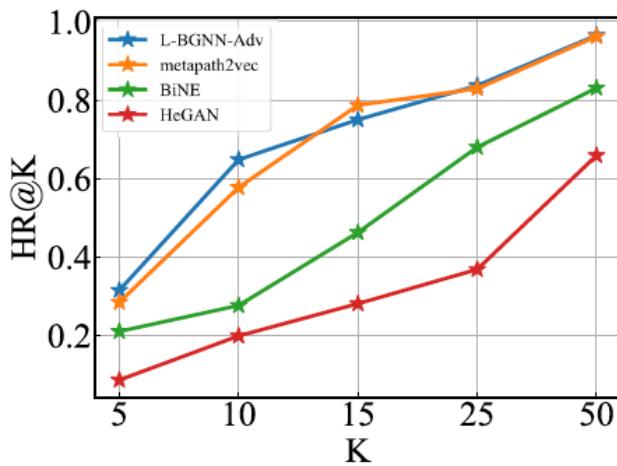
TABLE III
PERFORMANCE COMPARISON FOR THE LINK PREDICTION TASK

	DBLP			Yelp			Amazon		
Methods	Accuracy	Macro F1	AUC	Accuracy	Macro F1	AUC	Accuracy	Macro F1	AUC
node2vec	0.679	0.669	0.764	0.666	0.666	0.720	0.665	0.665	0.719
ANRL	0.791	0.790	0.853	0.840	0.839	0.890	0.717	0.717	0.795
VGAE	0.553	0.455	0.701	0.760	0.749	0.908	0.623	0.608	0.726
GraphSAGE-MEAN	0.736	0.735	0.760	0.707	0.707	0.712	0.702	0.702	0.704
HeGAN	0.571	0.569	0.610	0.516	0.516	0.520	0.519	0.518	0.525
FeatWalk	0.560	0.553	0.567	0.562	0.554	0.579	0.624	0.623	0.695
metapath2vec	0.822	0.820	0.903	0.846	0.846	0.923	0.741	0.741	0.815
HIN2Vec	0.855	0.854	0.926	0.752	0.744	0.877	0.678	0.672	0.767
BiNE	0.600	0.560	0.780	0.603	0.593	0.660	0.568	0.515	0.713
BiANE	0.858	0.858	0.927	0.857	0.856	0.918	0.739	0.736	0.820
L-BGNN-Adv	0.861	0.860	0.931	0.882	0.882	0.932	0.749	0.749	0.827



- Train a logistic classifier on the node embeddings;
- Metrics: accuracy; macro F1; AUC;
- SOTA performance;

Evaluation: Recommendation



- Leave-one-out method for evaluation;
- Metrics: hit ratio @ k; normalized discounted cumulative gain @ k;
- SOTA performance;

Evaluation: Embedding Visualization

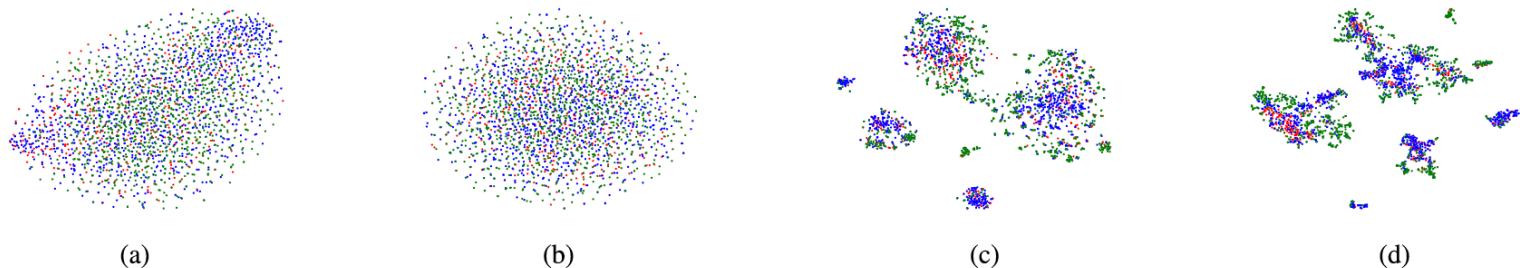


Fig. 6. Visualization of embedding results on the Yelp dataset, where colors are used to indicate different business categories. (a) BiNE. (b) HeGAN. (c) metapath2vec. (d) L-BGNN-Adv.

- t-SNE visualizes the node embeddings of the Yelp dataset;
- Different colors denote to different labels;
- Denser clusters and cleaner boundary;

Ablation Study: Effect of Intradomain Alignment (IDA)

TABLE IV

PERFORMANCE COMPARISON OF L-BGNN-ADV AGAINST SIMPLE DOMAIN ALIGNMENT, WHERE THE PERFORMANCE METRIC IS THE MICRO F1 SCORE

	Group	Cora	CiteSeer	PubMed
Raw features	0.501	0.693	0.663	0.754
Aggregated features	0.555	0.825	0.612	0.757
Concatenated features	0.556	0.797	0.637	0.770
L-BGNN-Adv	0.622	0.859	0.768	0.857

TABLE V

PERFORMANCE COMPARISON OF L-BGNN-ADV AND L-BGNN-MLP, WHERE THE PERFORMANCE METRIC IS THE MICRO F1 SCORE

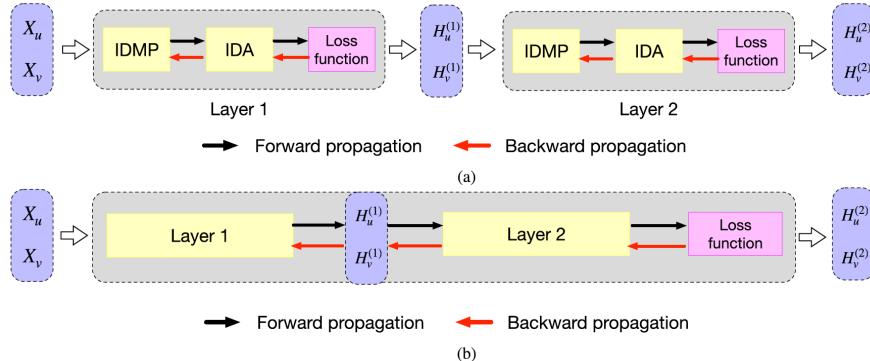
	Group	Cora	CiteSeer	PubMed
L-BGNN-MLP	0.582	0.784	0.756	0.846
L-BGNN-Adv	0.622	0.859	0.768	0.857

- Naive feature engineering;

- Linear domain alignment

$$\mathcal{L}_u = \left\| \text{MLP}\left(\mathbf{H}_{v \rightarrow u}^{(k)}\right) - \mathbf{H}_u^{(k)} \right\|_F^2$$

Ablation Study: Effect of Layerwise Training



- Adopt concatenation in hidden layers;

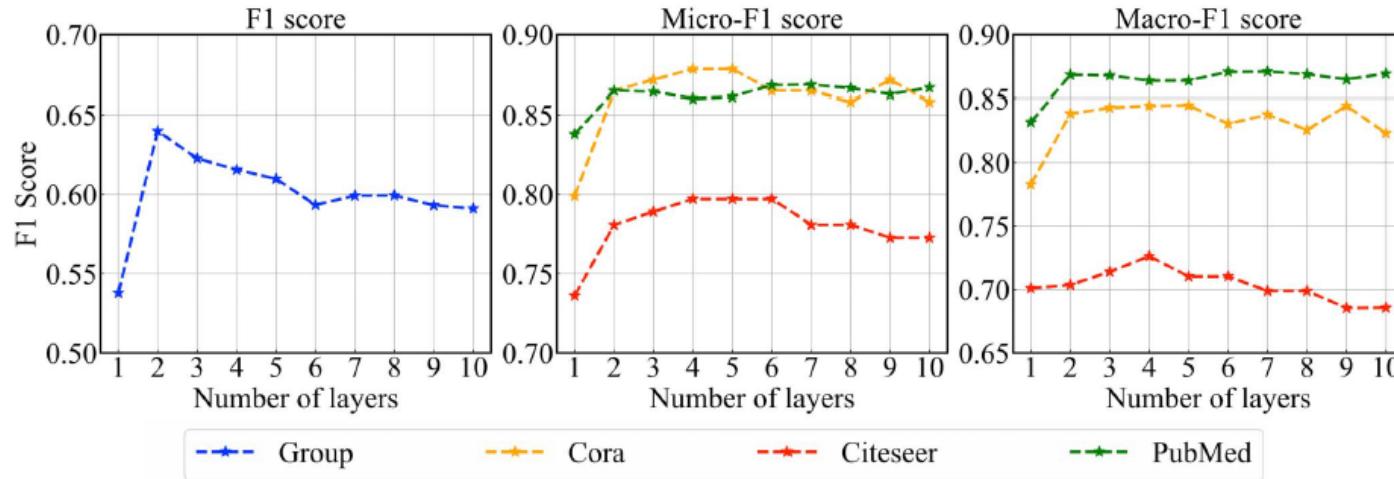
$$\begin{cases} \mathbf{H}_u^{(k+1)} = \text{CONCAT}\left(\mathbf{H}_{v \rightarrow u}^{(k)}, \mathbf{H}_u^{(k)}\right) \\ \mathbf{H}_v^{(k+1)} = \text{CONCAT}\left(\mathbf{H}_{u \rightarrow v}^{(k)}, \mathbf{H}_v^{(k)}\right). \end{cases}$$

- Same adversarial alignment in the end;

	End-to-end training	Layerwise training
Group	OOM	0.622
Cora	0.837/0.809	0.859/0.831
CiteSeer	0.685/0.642	0.768/0.698
PubMed	0.859 /0.859	0.857/ 0.860

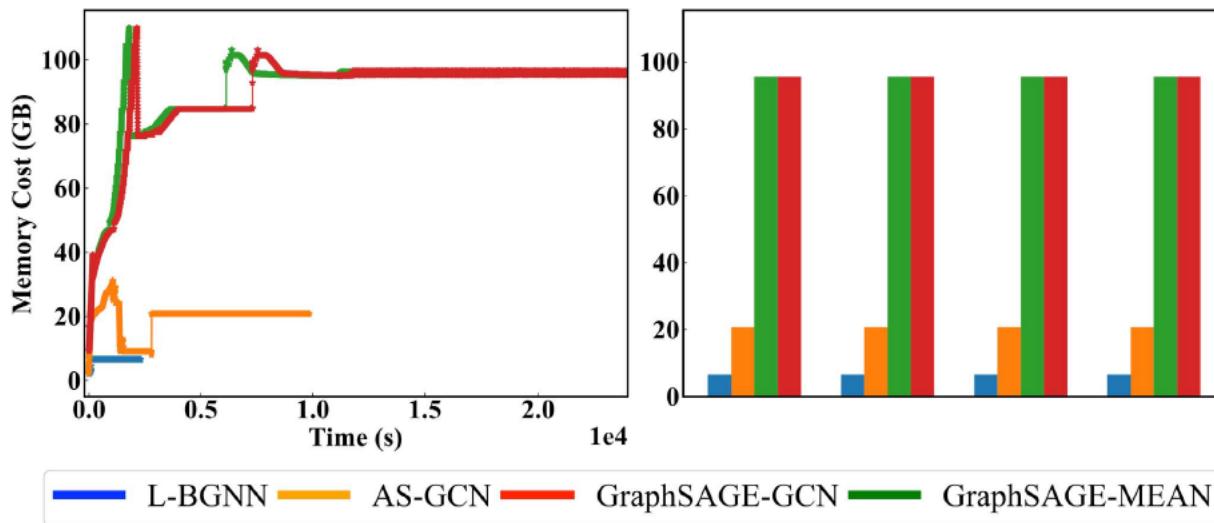
- End-to-end training fails on large-scale graphs;
- Layerwise training performs better;

Ablation Study: Effect of Layerwise Training



- Increase the number of training layers;
- Performance rises first, stabilizes next, and drops at the end;

Ablation Study: Efficiency for Large-scale Graphs



- Benchmarking methods: GNNs specific for large-scale graphs;
- Lower memory cost;
- Faster training speed;

Outline

- Background & Related Work
- Challenges for Bipartite Graphs
- Our Method: L-BGNN
- Experiments & Evaluation
- Conclusion

Conclusion

- Background & Related Work
 - Graph Representation Learning
 - Graph Neural Networks
- Challenges for Bipartite Graphs
 - Domain Inconsistency
 - Scalable & Label Efficiency
- Our Method: L-BGNN
 - Interdomain Message Passing (IDMP)
 - Intradomain Alignment (IDA)
 - Layerwise Training
- Experiments & Evaluation
 - SOTA Performance on Multiple Tasks
 - Ablation Study



Thank You!

Q&A